

# Final\_CC\_Default

February 21, 2021

## 1 Prediction of Credit Card Default for Taiwanese Customers

## 2 Contents

- Introduction
- Importing Packages
- Uploading Data
- Creating Train, Validation, and Testing Sets
- Data Cleaning
- Exploratory Data Analysis
- Feature Engineering

## 3 Introduction

This dataset contains information on default payments, demographic factors, credit data, history of payment, and bill statements of credit card customers in Taiwan from April 2005 to September 2005.

There are 25 variables:

- ID: ID of each client
- LIMIT\_BAL: Amount of given credit in NT dollars (includes individual and family/supplementary credit)
- SEX: Gender
  - 1=male
  - 2=female
- EDUCATION:
  - 1=graduate school
  - 2=university
  - 3=high school
  - 4=others
  - 5=unknown
  - 6=unknown
- MARRIAGE: Marital status
  - 1=married

- 2=single
  - 3=others)
- AGE: Age in years
- PAY\_0: Repayment status in September, 2005
  - -1=pay duly
  - 1=payment delay for one month
  - 2=payment delay for two months ....
  - 8=payment delay for eight months
  - 9=payment delay for nine months and above
- PAY\_2: Repayment status in August, 2005 (scale same as above)
- PAY\_3: Repayment status in July, 2005 (scale same as above)
- PAY\_4: Repayment status in June, 2005 (scale same as above)
- PAY\_5: Repayment status in May, 2005 (scale same as above)
- PAY\_6: Repayment status in April, 2005 (scale same as above)
- BILL\_AMT1: Amount of bill statement in September,2005 (NT dollar)
- BILL\_AMT2: Amount of bill statement in August, 2005 (NT dollar)
- BILL\_AMT3: Amount of bill statement in July, 2005 (NT dollar)
- BILL\_AMT4: Amount of bill statement in June, 2005 (NT dollar)
- BILL\_AMT5: Amount of bill statement in May, 2005 (NT dollar)
- BILL\_AMT6: Amount of bill statement in April, 2005 (NT dollar)
- PAY\_AMT1: Amount of previous payment in September, 2005 (NT dollar)
- PAY\_AMT2: Amount of previous payment in August, 2005 (NT dollar)
- PAY\_AMT3: Amount of previous payment in July, 2005 (NT dollar)
- PAY\_AMT4: Amount of previous payment in June, 2005 (NT dollar)
- PAY\_AMT5: Amount of previous payment in May, 2005 (NT dollar)
- PAY\_AMT6: Amount of previous payment in April, 2005 (NT dollar)
- default.payment.next.month: Default payment
  - 1=yes
  - 0=no

## 4 Importing Packages

```
[2]: # Importing Packages
import numpy as np
import pandas as pd
import re
import json
import requests
import warnings
warnings.filterwarnings('ignore')
pd.set_option("display.max_rows", 999)
pd.set_option("display.max_columns", 999)
from collections import Counter
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
%matplotlib inline
```

```

plt.style.use("fivethirtyeight")
import pickle

from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier,
    ↳ RandomForestClassifier, BaggingClassifier
from sklearn.model_selection import train_test_split, KFold, cross_val_score,
    ↳ cross_val_predict, GridSearchCV
from sklearn.feature_selection import RFECV
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score,
    ↳ recall_score, f1_score, roc_auc_score,
    ↳ classification_report, balanced_accuracy_score
from sklearn.preprocessing import StandardScaler
import xgboost as xgb
from xgboost.sklearn import XGBClassifier

%reload_ext autoreload
%autoreload 2
from utils import *

```

```

[3]: df = pd.read_excel("data/default of credit card clients.xls")
new_header = df.iloc[0]
df = df[1:]
df.columns = new_header
df = df.rename(columns={"default payment next month": "default"})

```

## 5 Create Dataset Splits

```

[4]: X = df.drop(["default"], axis=1)
y = df["default"]
X_train, X_val, y_train, y_val = train_test_split(X, y, train_size=0.8,
    ↳ random_state=42)
X_tr, X_tt, y_tr, y_tt = train_test_split(X_train, y_train, train_size=0.875,
    ↳ random_state=42)
train = pd.concat([X_tr, y_tr], axis=1)
val = pd.concat([X_val, y_val], axis=1)
tr = train.drop(["ID"], axis=1)
val = val.drop(["ID"], axis=1)

```

## 6 Data Cleaning

- The data cleaning process involved coonverting Taiwanese dollars to American dollars to facilitate understanding of the numbers involved and then converting them into integers.
- There were anomalous values in the marriage category, so observations with the value of 0 were converted into 3, which represented other.
- There were anomalous values in the education category, so observations with the value of 0, 5, or 6 were all lumped into the 4 or other category.
- In the behind1 - behind6 categories, I was originally going to convert all the observations with the negative values into 0, but since there were so many observations with -1 and -2, it couldn't have been an anomaly or mistake.

```
[5]: url = 'https://openexchangerates.org/api/latest.json?
      ↪app_id=c51b1508fb4145259b1c2fade72a2c04'
      response = requests.get(url)
      data = response.json()
      rate = data['rates']['TWD']
```

```
[6]: data = [tr, val]
      for d in data:
          d.rename(columns={"PAY_0": "behind1", "PAY_2": "behind2", "PAY_3": "behind3", "PAY_4": "behind4", "PAY_5": "behind5", "PAY_6": "behind6", "BILL_AMT1": "billed1", "BILL_AMT2": "billed2", "BILL_AMT3": "billed3", "BILL_AMT4": "billed4", "BILL_AMT5": "billed5", "BILL_AMT6": "billed6", "PAY_AMT1": "paid1", "PAY_AMT2": "paid2", "PAY_AMT3": "paid3", "PAY_AMT4": "paid4", "PAY_AMT5": "paid5", "PAY_AMT6": "paid6", "SEX": "gender", "EDUCATION": "education", "MARRIAGE": "marriage", "AGE": "age", "LIMIT_BAL": "limit"}, inplace=True)
          d[['limit']] = d[['limit']]/rate
          d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']] = d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']].divide(rate, axis=1).astype(int)
          d[['paid1', 'paid2', 'paid3', 'paid4', 'paid5', 'paid6']] = d[['paid1', 'paid2', 'paid3', 'paid4', 'paid5', 'paid6']].divide(rate, axis=1).astype(int)
          d['limit'] = d['limit'].apply(lambda x: round(x, 2))
          d.replace({'marriage': {0:3}}, inplace=True)
          d.replace({'education': {5:4, 0:4, 6:4}}, inplace=True)
```

```
[7]: tr.head()
```

```
[7]:      limit  gender  education  marriage  age  behind1  behind2  behind3  \
6191  1788.52      2          2          1   44         0         0         0
16054  5723.26      2          3          1   46        -1        -1        -1
19706  3577.04      2          2          1   47        -1        -1        -1
23128  6080.96      2          2          1   29         0         0         0
28516  5365.56      2          1          2   33        -2        -2        -2
```

	behind4	behind5	behind6	billed1	billed2	billed3	billed4	billed5	\
6191	0	0	0	1630	1498	1277	799	846	
16054	0	-1	-1	890	83	173	147	142	
19706	-1	-1	-2	238	238	0	224	-14	
23128	0	0	0	2829	2238	2264	2285	1556	
28516	-2	-2	-2	872	960	1169	1196	994	

	billed6	paid1	paid2	paid3	paid4	paid5	paid6	default
6191	980	107	178	107	107	178	33	0
16054	30	83	173	35	142	30	941	0
19706	-14	238	0	224	0	0	0	1
23128	1573	79	89	92	60	67	75	0
28516	80	966	1170	1197	994	80	6061	0

```
[8]: tr.describe()
```

```
[8]:
```

	limit	education	marriage	billed1	billed2	\
count	21000.000000	21000.000000	21000.000000	21000.000000	21000.000000	
mean	5981.340686	1.842810	1.555333	1827.979905	1758.19519	
std	4634.450783	0.746378	0.522538	2628.013848	2545.94903	
min	357.700000	1.000000	1.000000	-5922.000000	-2495.00000	
25%	1788.520000	1.000000	1.000000	127.000000	108.00000	
50%	5007.850000	2.000000	2.000000	803.000000	766.00000	
75%	8584.890000	2.000000	2.000000	2387.250000	2263.50000	
max	35770.370000	4.000000	3.000000	34500.000000	35195.00000	

	billed3	billed4	billed5	billed6	paid1	\
count	21000.000000	21000.000000	21000.000000	21000.000000	21000.000000	
mean	1679.613952	1543.949857	1443.855667	1393.781810	204.618429	
std	2489.800493	2299.447009	2182.100061	2134.592138	626.349241	
min	-5625.000000	-6080.000000	-2909.000000	-7477.000000	0.000000	
25%	98.000000	84.000000	63.000000	46.000000	35.000000	
50%	717.500000	680.500000	647.000000	612.000000	75.000000	
75%	2135.000000	1943.250000	1798.000000	1767.000000	179.000000	
max	59525.000000	31892.000000	33165.000000	34399.000000	31247.000000	

	paid2	paid3	paid4	paid5	paid6
count	21000.000000	21000.000000	21000.000000	21000.000000	21000.000000
mean	214.418048	188.432952	175.843714	172.090667	184.061095
std	897.321096	667.001026	601.041739	558.702990	631.634456
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	29.000000	13.000000	10.000000	8.000000	4.000000
50%	71.000000	64.000000	53.000000	53.000000	53.000000
75%	178.000000	160.000000	143.000000	144.000000	143.000000
max	60246.000000	32051.000000	22213.000000	14951.000000	18856.000000

**\*\* Observations: \*\***

- No missing data
- There were anomalous values for education and marriage, and the anomalous values were reassigned under other.
- Did not reassign -2 and -1 to 0 for 'behind' features despite being anomalous because they were so many -2 and -1. There must be so significance to those values.

## 7 Exploratory Data Anlaysis

```
[9]: # organize features into categorical and continuous
categorical = tr[['gender', 'marriage', 'education', 'behind1', 'behind2', 'behind3', 'behind4', 'behind5', 'behind6']]
continuous = tr[['limit', 'age', 'billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6', 'paid1', 'paid2', 'paid3', 'paid4', 'paid5', 'paid6']]
cat_col = categorical.columns
cont_col = continuous.columns
```

```
[10]: # display distributions of all the continuous variables

# con_1 = pd.melt(tr, value_vars = cont_col)
# sns.set_theme(style="darkgrid", font='serif', context='talk')
# g = sns.FacetGrid(con_1, col='variable', col_wrap=3, sharex=False, sharey=False, height=4)
# g = g.map(sns.distplot, 'value', color='r')
# g.set_xticklabels(rotation=45)
# g.fig.subplots_adjust(top=0.9)
# g.fig.suptitle("Distributions of Continuous Features")
# g.fig.tight_layout()
# plt.savefig("../images/distplot.png")
```

**\*\* Observations: \*\***

- It is hard to observe any trends with the paid features.

```
[11]: # Use bar graphs of the distribution of data for categorical variables

# cat_1 = pd.melt(tr, value_vars=cat_col)
# sns.set_theme(style="darkgrid", font='serif', context='talk')
# g = sns.FacetGrid(cat_1, col='variable', col_wrap=3, sharex=False, sharey=False, height=4)
# g = g.map(sns.countplot, 'value', color='dodgerblue')
# g.set_xticklabels()
# g.fig.subplots_adjust(top=0.9)
# g.fig.suptitle("Distributions of Categorical Features")
# g.fig.tight_layout()
# plt.savefig("../images/countplot.png")
```

```
[12]: yes = tr.default.sum()
no = len(tr)-yes
perc_y = round(yes/len(tr)*100, 1)
perc_n = round(no/len(tr)*100, 1)

# plt.figure(figsize=(8,6))
# sns.set_theme(style="darkgrid", font='serif', context='talk')
# sns.countplot('default', data=tr)
# plt.title('Credit Card Baseline Default', size=16)
# plt.box(False);
# plt.savefig("../images/baseline.png")
```

- There is class imbalance in the dataset. Our baseline indicates

```
[14]: print("Number of Total Non-Defaulters: ", yes)
print("Number of Defaulters: ", no)
print("Percentage of Non-Defaulters: ", perc_y)
print("Percentage of Defaulters: ", perc_n)

pd.DataFrame
default = pd.DataFrame(data = {"Training Dataset": [yes, no, perc_y, perc_n]},
                        index = ["Number of Total Non-Defaulters: ", "Number of_
↳Defaulters: ", "Percentage of Non-Defaulters: ", "Percentage of Defaulters:
↳"])
default
```

Number of Total Non-Defaulters: 4656

Number of Defaulters: 16344

Percentage of Non-Defaulters: 22.2

Percentage of Defaulters: 77.8

```
[14]: Training Dataset
Number of Total Non-Defaulters: 4656.0
Number of Defaulters: 16344.0
Percentage of Non-Defaulters: 22.2
Percentage of Defaulters: 77.8
```

```
[15]: # subset = tr[['gender', 'education', 'marriage', 'behind1', 'behind2',
↳'behind3', 'behind4', 'behind5', 'behind6', 'default']]
# f, axes = plt.subplots(3, 3, figsize=(15, 12), facecolor='white')
# sns.set_theme(style="darkgrid", font='serif', context='paper')
# f.suptitle('Frequency of Categorical Variables', size=16)
# ax1 = sns.countplot(x="gender", hue="default", data=subset, ax=axes[0,0])
# ax2 = sns.countplot(x="education", hue="default", data=subset, ax=axes[0,1])
# ax3 = sns.countplot(x="marriage", hue="default", data=subset, ax=axes[0,2])
# ax4 = sns.countplot(x="behind1", hue="default", data=subset, ax=axes[1,0])
# ax5 = sns.countplot(x="behind2", hue="default", data=subset, ax=axes[1,1])
# ax6 = sns.countplot(x="behind3", hue="default", data=subset, ax=axes[1,2])
```

```
# ax7 = sns.countplot(x="behind4", hue="default", data=subset, ax=axes[2,0])
# ax8 = sns.countplot(x="behind5", hue="default", data=subset, ax=axes[2,1])
# ax9 = sns.countplot(x="behind6", hue="default", data=subset, ax=axes[2,2])
# plt.savefig("../images/default_freq_by_cat.png")
```

**\*\* Observations: \*\***

- **gender**, **education**, and **marriage** doesn't seem to change with each group in terms of proportions. **Behind** seems to have some correlation with **default**. That would make sense since being behind in payments would make it more likely that you would default next month.
- There isn't a very clear distinction between the distribution of **default** on any of the demographic data. However, you do see quite some distribution differences of target classes for monthly repayment status (**behind1-behind6**).

```
[16]: education = tr.groupby(['education', 'default']).size().unstack(1)
education
# education.plot(kind="bar", stacked=True)
# plt.title("Distribution Count of Educational Level and Default Status",
↳ size=14)
# plt.savefig("../data/stacked_bar2.png")
```

```
[16]: default      0      1
education
1          6013   1424
2          7408   2341
3          2626    866
4           297     25
```

**\*\* No clear relationship with education. The proportion doesn't seem to change with each group. \*\***

```
[17]: marriage = tr.groupby(['marriage', 'default']).size().unstack(1)
marriage
# marriage.plot(kind="bar", stacked=True)
# plt.title("Distribution of Default Status for Marital Status", size=14)
# plt.savefig("../images/stacked_bar3.png")
```

```
[17]: default      0      1
marriage
1          7354   2258
2          8778   2336
3           212     62
```



## 8 Vanilla Model

```
[ ]: logreg = LogisticRegression(solver="liblinear", random_state=42).fit(X_tr_dum, y_tr)
y_pred_log_tr = logreg.predict(X_tr_dum)
y_pred_log_val = logreg.predict(X_val_dum)

rfc = RandomForestClassifier().fit(X_tr, y_tr)
y_pred_rfc_tr = rfc.predict(X_tr)
y_pred_rfc_val = rfc.predict(X_val)

dtc = DecisionTreeClassifier().fit(X_tr, y_tr)
y_pred_dtc_tr = dtc.predict(X_tr)
y_pred_dtc_val = dtc.predict(X_val)

knn = KNeighborsClassifier().fit(X_tr, y_tr)
y_pred_knn_tr = knn.predict(X_tr)
y_pred_knn_val = knn.predict(X_val)

gnb = GaussianNB().fit(X_tr, y_tr)
y_pred_gnb_tr = gnb.predict(X_tr)
y_pred_gnb_val = gnb.predict(X_val)

lda = LinearDiscriminantAnalysis().fit(X_tr, y_tr)
y_pred_lda_tr = lda.predict(X_tr)
y_pred_lda_val = lda.predict(X_val)

abc = AdaBoostClassifier().fit(X_tr, y_tr)
y_pred_abc_tr = abc.predict(X_tr)
y_pred_abc_val = abc.predict(X_val)

gbc = GradientBoostingClassifier().fit(X_tr, y_tr)
y_pred_gbc_tr = gbc.predict(X_tr)
y_pred_gbc_val = gbc.predict(X_val)

xgb = XGBClassifier().fit(X_tr, y_tr)
y_pred_xgb_tr = xgb.predict(X_tr)
y_pred_xgb_val = xgb.predict(X_val)
```

```
[18]: baseline = pd.read_csv("data/charts/baseline.csv")
baseline
```

```
[18]:
```

	Unnamed: 0	Accuracy	F1 Score	Recall	Precision	\
0	Logistic Regression	0.811500	0.360656	0.242955	0.699561	
1	Random Forest Classifier	0.814833	0.461464	0.362529	0.634667	
2	Decision Tree Classifier	0.730667	0.399257	0.408987	0.389978	
3	K-Nearest Neighbors	0.798000	0.447080	0.373191	0.557452	

4	Gaussian Naive Bayes	0.724000	0.498486	0.626809	0.413776
5	Linear Discriminant Analysis	0.810333	0.367778	0.252094	0.679671
6	AdaBoost Classifier	0.815667	0.425753	0.312262	0.668842
7	Gradient Boosting Classifier	0.820833	0.468085	0.360244	0.668079
8	XGBoost Classifier	0.816833	0.469338	0.370145	0.641161

```

PR AUC
0 0.486829
1 0.513445
2 0.289090
3 0.416605
4 0.480981
5 0.480476
6 0.523430
7 0.545925
8 0.518488

```

## 9 Feature Engineering

- `age_bin`: 1 = young adult, 2 = middle age, 3 = senior
- `gen-mar`: interaction between gender and marriage status
- `gen-age`: interaction between age and gender
- `'avail...'`: fraction of estimated available balance based on what is billed per month

```

[19]: data = [tr, val]

# create features for demographic variables
for d in data:
    d['age_bin'] = 0
    d.loc[((d['age'] > 20) & (d['age'] < 30)) , 'age_bin'] = 1
    d.loc[((d['age'] >= 30) & (d['age'] < 60)) , 'age_bin'] = 2
    d.loc[((d['age'] >= 60) & (d['age'] < 81)) , 'age_bin'] = 3
    # create categories for single, married, divorced males and females
    d['gen-mar'] = d['gender'] + d['marriage']
    # create categories for young, middle age and senior males and females
    d['gen-age'] = d['gender'] + d['age_bin']

# feature for credit use percentage: fraction of estimated available balance
→based on what is billed per month
# (credit limit - monthly billed amount) / credit limit
for d in data:
    d['avail6'] = (d.limit - d.billed6) / d.limit
    d['avail5'] = (d.limit - d.billed5) / d.limit
    d['avail4'] = (d.limit - d.billed4) / d.limit
    d['avail3'] = (d.limit - d.billed3) / d.limit
    d['avail2'] = (d.limit - d.billed2) / d.limit

```

```

    d['avail1'] = (d.limit - d.billed1) / d.limit
    d['avg_av'] = (d.avail1 + d.avail2 + d.avail3 + d.avail4 + d.avail5 + d.
↳avail6) / 6

# create a feature that indicates whether a client has had a delayed payment or
↳not
def delayed_payment(d):
    if (d.behind1 > 0) or (d.behind2 > 0) or (d.behind3 > 0) or (d.behind4 > 0)
↳or (d.behind5 > 0) or (d.behind6 > 0):
        return 1
    else:
        return 0
for d in data:
    d['delayed'] = d.apply(delayed_payment, axis=1)

# create feature for the total number of months with delayed payment status for
↳a particular client
def total_months_with_delayed_payments(d):
    count = 0
    if (d.behind1 > 0):
        count += 1
    if (d.behind2 > 0):
        count += 1
    if (d.behind3 > 0):
        count += 1
    if (d.behind4 > 0):
        count += 1
    if (d.behind5 > 0):
        count += 1
    if (d.behind6 > 0):
        count += 1
    return count
for d in data:
    d['latemths'] = d.apply(total_months_with_delayed_payments, axis=1)

# the ratio of amount paid and amount billed
for d in data:
    d['pperb1'] = d.paid1 / d.billed2
    d['pperb2'] = d.paid2 / d.billed3
    d['pperb3'] = d.paid3 / d.billed4
    d['pperb4'] = d.paid4 / d.billed5
    d['pperb5'] = d.paid5 / d.billed6

# remove any infinity and NaN values
datasets = ['pperb1', 'pperb2', 'pperb3', 'pperb4', 'pperb5']
for data in datasets:
    tr.replace({data: {np.inf: 0, np.nan: 0}}, inplace=True)

```

```
val.replace({data: {np.inf: 0, np.nan: 0}}, inplace=True)
```

```
[20]: # plt.style.use("fivethirtyeight")
# sns.set_theme(style="darkgrid", font='serif', context='paper')
# plt.figure(figsize = (20,16))
# plt.title('Pearson Correlation of Features', y = 1.05, size = 20)
# g = sns.heatmap(tr.corr(), cmap='RdBu', square=True, linecolor='white',
→linewidths=0.2)
# plt.savefig("../images/correlation_matrix_2.png")
```

\*\* This includes my engineered features. Default seems to be correlated with two of my engineered features, delayed and latemnth. Delayed is whether you have had a delayed payment during the 6 month history or not. latemnth is the total number of months you were given a status of behind in payments. Seems to be correlated with behind1 and limit.\*\*

```
[22]: pickle_in = open("data/pickles/training_features.pickle","rb")
train2 = pickle.load(pickle_in)
pickle_in = open("data/pickles/validate_features.pickle","rb")
validate2 = pickle.load(pickle_in)
```

```
[23]: X_train2 = train2.drop(["default"], axis=1)
y_tr = train2["default"]
X_validate2 = validate2.drop(["default"], axis=1)
y_val = validate2["default"]
```

```
[24]: # # Grab indices of columns for creating dummy variables and create dataframe
→with dummy variables
dum_feat = X_train2[['gender', 'education', 'marriage', 'age_bin', 'gen-mar',
→'gen-age']]
dum_index = dum_feat.columns
tr_dum = pd.get_dummies(data=dum_feat, columns=dum_index, drop_first=True,
→prefix=['sex', 'edu', 'mar', 'agebin', 'sexmar', 'sexage'])
cont_feat = X_train2.drop(['gender', 'education', 'marriage', 'age_bin',
→'gen-mar', 'gen-age'], axis=1)
X_train2_dum = cont_feat.join(tr_dum)
X_train2_dum.head()
```

```
[24]:
```

	limit	age	behind1	behind2	behind3	behind4	behind5	behind6	\
0	1790.26	44	0	0	0	0	0	0	
1	5728.83	46	-1	-1	-1	0	-1	-1	
2	3580.52	47	-1	-1	-1	-1	-1	-2	
3	6086.88	29	0	0	0	0	0	0	
4	5370.78	33	-2	-2	-2	-2	-2	-2	

  

	billed1	billed2	billed3	billed4	billed5	billed6	paid1	paid2	\
0	1631.93	1500.45	1278.35	800.60	847.12	981.81	107.99	179.13	
1	891.69	83.71	173.87	147.77	143.04	30.15	83.89	173.87	

2	238.68	238.68	0.00	224.50	-14.18	-14.18	238.68	0.00
3	2831.87	2240.51	2267.08	2288.06	1557.71	1575.25	80.02	89.26
4	873.40	961.26	1170.90	1198.01	995.38	80.96	966.99	1171.37

  

	paid3	paid4	paid5	paid6	avail6	avail5	avail4	avail3	\
0	107.42	107.42	179.03	33.08	0.451582	0.526817	0.552802	0.285942	
1	35.81	143.04	30.15	942.14	0.994737	0.975032	0.974206	0.969650	
2	224.50	0.00	0.00	0.00	1.003960	1.003960	0.937300	1.000000	
3	92.56	60.26	68.07	75.58	0.741206	0.744087	0.624100	0.627546	
4	1198.58	995.67	80.96	6067.73	0.984926	0.814668	0.776939	0.781987	

  

	avail2	avail1	avg_av	delayed	latemths	pperb1	pperb2	\
0	0.161882	0.088440	0.344578	0	0	0.071972	0.140126	
1	0.985388	0.844350	0.957227	0	0	1.002150	1.000000	
2	0.933339	0.933339	0.968650	0	0	1.000000	0.000000	
3	0.631912	0.534758	0.650602	0	0	0.035715	0.039372	
4	0.821020	0.837379	0.836153	0	0	1.005961	1.000401	

  

	pperb3	pperb4	pperb5	sex_2	edu_2	edu_3	edu_4	mar_2	mar_3	\
0	0.134174	0.126806	0.182347	1	1	0	0	0	0	
1	0.242336	1.000000	1.000000	1	0	1	0	0	0	
2	1.000000	-0.000000	-0.000000	1	1	0	0	0	0	
3	0.040453	0.038685	0.043212	1	1	0	0	0	0	
4	1.000476	1.000291	1.000000	1	0	0	0	1	0	

  

	agebin_2	agebin_3	sexmar_3	sexmar_4	sexmar_5	sexage_3	sexage_4	\
0	1	0	1	0	0	0	1	
1	1	0	1	0	0	0	1	
2	1	0	1	0	0	0	1	
3	0	0	1	0	0	1	0	
4	1	0	0	1	0	0	1	

  

	sexage_5
0	0
1	0
2	0
3	0
4	0

```
[25]: dum_feat2 = X_validate2[['gender', 'education', 'marriage', 'age_bin',
    ↪ 'gen-mar', 'gen-age']]
dum_index2 = dum_feat2.columns
val_dum = pd.get_dummies(data=dum_feat2, columns=dum_index2, drop_first=True,
    ↪ prefix=['sex', 'edu', 'mar', 'agebin', 'sexmar', 'sexage'])
cont_feat2 = X_validate2.drop(['gender', 'education', 'marriage', 'age_bin',
    ↪ 'gen-mar', 'gen-age'], axis=1)
X_validate2_dum = cont_feat2.join(val_dum)
```

```
X_validate2_dum.head()
```

```
[25]:
```

	limit	age	behind1	behind2	behind3	behind4	behind5	behind6	\
0	1074.16	25	0	0	0	0	0	0	
1	5370.78	26	0	0	0	0	0	0	
2	2506.36	32	0	0	0	0	0	0	
3	4654.68	49	0	0	0	0	0	-1	
4	1790.26	36	0	0	0	0	0	2	

  

	billed1	billed2	billed3	billed4	billed5	billed6	paid1	paid2	\
0	317.38	360.27	414.66	450.43	491.10	530.92	53.71	71.61	
1	4895.86	4498.96	4177.89	3637.13	2783.53	2766.45	160.62	151.64	
2	2510.73	2473.42	2453.73	2497.52	2510.34	2513.96	87.04	111.43	
3	740.38	678.72	579.04	605.04	402.31	248.63	57.65	64.74	
4	3373.85	1705.58	1516.74	700.85	726.67	696.02	71.61	53.71	

  

	paid3	paid4	paid5	paid6	avail6	avail5	avail4	avail3	\
0	53.71	53.71	53.71	71.61	0.505735	0.542806	0.580668	0.613968	
1	113.18	94.78	95.56	95.56	0.484907	0.481727	0.322793	0.222107	
2	107.42	87.29	89.51	91.45	-0.003032	-0.001588	0.003527	0.020999	
3	251.14	0.97	251.03	157.83	0.946585	0.913569	0.870015	0.875600	
4	35.81	64.45	0.00	35.81	0.611218	0.594098	0.608521	0.152782	

  

	avail2	avail1	avg_av	delayed	latemths	pperb1	pperb2	\
0	0.664603	0.704532	0.602052	0	0	0.149083	0.172696	
1	0.162327	0.088427	0.293715	0	0	0.035702	0.036296	
2	0.013143	-0.001744	0.005217	0	0	0.035190	0.045412	
3	0.854185	0.840939	0.883482	0	0	0.084939	0.111806	
4	0.047300	-0.884559	0.188227	1	1	0.041986	0.035411	

  

	pperb3	pperb4	pperb5	sex_2	edu_2	edu_3	edu_4	mar_2	mar_3	\
0	0.119242	0.109367	0.101164	0	1	0	0	1	0	
1	0.031118	0.034050	0.034542	1	0	0	0	1	0	
2	0.043011	0.034772	0.035605	1	0	1	0	0	0	
3	0.415080	0.002411	1.009653	0	0	1	0	1	0	
4	0.051095	0.088692	0.000000	1	1	0	0	1	0	

  

	agebin_2	agebin_3	sexmar_3	sexmar_4	sexmar_5	sexage_3	sexage_4	\
0	0	0	1	0	0	0	0	
1	0	0	0	1	0	1	0	
2	1	0	1	0	0	0	1	
3	1	0	1	0	0	1	0	
4	1	0	0	1	0	0	1	

  

	sexage_5
0	0
1	0

```
2      0
3      0
4      0
```

```
[22]: scaler = StandardScaler().fit(X_train2_dum)
X_tr2_dum = scaler.transform(X_train2_dum)
X_val2_dum = scaler.transform(X_validate2_dum)
```

```
[23]: scaler2 = StandardScaler().fit(X_train2)
X_tr2 = scaler2.transform(X_train2)
X_val2 = scaler2.transform(X_validate2)
```

```
[ ]: logreg2 = LogisticRegression(solver="liblinear", random_state=42).
    ↪fit(X_tr2_dum, y_tr)
y_pred_log_tr2 = logreg2.predict(X_tr2_dum)
y_pred_log_val2 = logreg2.predict(X_val2_dum)

rfc2 = RandomForestClassifier().fit(X_tr2, y_tr)
y_pred_rfc_tr2 = rfc2.predict(X_tr2)
y_pred_rfc_val2 = rfc2.predict(X_val2)

dtc2 = DecisionTreeClassifier().fit(X_tr2, y_tr)
y_pred_dtc_tr2 = dtc2.predict(X_tr2)
y_pred_dtc_val2 = dtc2.predict(X_val2)

knn2 = KNeighborsClassifier().fit(X_tr2, y_tr)
y_pred_knn_tr2 = knn2.predict(X_tr2)
y_pred_knn_val2 = knn2.predict(X_val2)

gnb2 = GaussianNB().fit(X_tr2, y_tr)
y_pred_gnb_tr2 = gnb2.predict(X_tr2)
y_pred_gnb_val2 = gnb2.predict(X_val2)

lda2 = LinearDiscriminantAnalysis().fit(X_tr2, y_tr)
y_pred_lda_tr2 = lda2.predict(X_tr2)
y_pred_lda_val2 = lda2.predict(X_val2)

abc2 = AdaBoostClassifier().fit(X_tr2, y_tr)
y_pred_abc_tr2 = abc2.predict(X_tr2)
y_pred_abc_val2 = abc2.predict(X_val2)

gbc2 = GradientBoostingClassifier().fit(X_tr2, y_tr)
y_pred_gbc_tr2 = gbc2.predict(X_tr2)
y_pred_gbc_val2 = gbc2.predict(X_val2)

xgb2 = XGBClassifier().fit(X_tr2, y_tr)
y_pred_xgb_tr2 = xgb2.predict(X_tr2)
```

```
y_pred_xgb_val2 = xgb2.predict(X_val2)
```

```
[27]: features_model = pd.read_csv("data/charts/features_model.csv")
      features_model
```

```
[27]:
```

		Unnamed: 0	Accuracy	F1 Score	Recall	Precision	\
0	Logistic Regression	2	0.809000	0.399371	0.290175	0.640336	
1	Random Forest Classifier	2	0.816500	0.467344	0.367860	0.640584	
2	Decision Tree Classifier	2	0.725167	0.393527	0.407464	0.380512	
3	K-Nearest Neighbors	2	0.789500	0.425125	0.355674	0.528281	
4	Gaussian Naive Bayes	2	0.278167	0.374042	0.985529	0.230824	
5	Linear Discriminant Analysis	2	0.809000	0.437684	0.339680	0.615172	
6	AdaBoost Classifier	2	0.818500	0.451385	0.341203	0.666667	
7	Gradient Boosting Classifier	2	0.820500	0.463378	0.354151	0.670029	
8	XGBoost Classifier	2	0.813333	0.458937	0.361767	0.627477	

  

	PR AUC
0	0.500013
1	0.517269
2	0.284835
3	0.405541
4	0.476403
5	0.500475
6	0.525636
7	0.544927
8	0.529126

**\*\* Observations: \*\***

In imbalanced datasets, we don't want to use accuracy as our gold standard metric since it is easy to get a high accuracy score by simply classifying all observations as the majority class.

ROC AUC score is equivalent to calculating the rank correlation between predictions and targets, i.e. how good at ranking predictions the model is. It tells us what is the probability that a positive instance randomly chosen is ranked higher than a negative instance randomly chosen. Generally, the ROC AUC is not used for imbalanced datasets because the FPR for highly imbalanced datasets is pulled down to a large number of true negatives.

With the F1 score, it combines both recall and precision into one metric by calculating the harmonic mean. It is worth noting that the F1 score is a special case of the F-beta score, where 1 indicates that we care about recall and precision equally. For a F2 score, we care about more recall more than precision, in fact, twice as much. Both F1 score and accuracy is calculated on the predicted classes not the prediction scores. We can adjust the threshold to finetune the F1 score, but accuracy also depends on the threshold.

PR AUC score shows the tradeoff between precision and recall at every threshold, where a high score or area represents both high recall and precision. High precision relates to a low FPR, and high recall relates to a low FNR. A high score means that the classifier is returning accurate results (high precision), as well as returning a majority of all positive results (high recall).



- To be expected, Random Forest and Decision Tree is overfit with high training accuracy and much lower validation accuracy.
- Tree-based and ensemble classifiers have the most potential with the highest PR AUC scores: Random Forest, AdaBoost, Gradient Boosting, XGBoost
- Decision Tree once again has an exceptionally low PR AUC score
- Gaussian Naive Bayes has an exceptionally low accuracy score in the second model
- A change in the F1 score between first and second model is accompanied by changes in recall and precision, which is to be expected since there is a tradeoff between recall and precision.
- Gaussian Bayes Classifier metrics changes significantly between the two models
- For both models, Gradient Boosting, AdaBoost, XGBoost, Random Forest have the highest PR AUC scores

## 10 Feature Selection

**\*\* Observations: \*\***

Top Features:

- Decision Tree Top 10: behind1, age, latemnth, avail1, avail2, billed, avail5, limit, paid1, avg\_av....
- Random Forest Top 10: behind1, age, latemnth, limit, avg\_av, avail1, billed1, delayed, avail2, behind2....
- XGBoost FI: delayed, behind1. (also latemnth, behind2)
- RFECV: limit, behind1, paid2, delayed, latemnth

Features to Remove:

- age\_bin, gender, marriage, gen-age, gen-mar, behind6, behind5, behind4, education

## 11 New Baseline Model

```
[29]: pickle_in = open("data/pickles/training_model.pickle", "rb")
train3 = pickle.load(pickle_in)
pickle_in = open("data/pickles/validate_model.pickle", "rb")
validate3 = pickle.load(pickle_in)
```

```
[30]: X_train3 = train3.drop(["default"], axis=1)
y_tr = train3["default"]
X_validate3 = validate3.drop(["default"], axis=1)
y_val = validate3["default"]

scaler3 = StandardScaler().fit(X_train3)
X_tr3 = scaler3.transform(X_train3)
X_val3 = scaler3.transform(X_validate3)
```

```
[ ]: logreg3 = LogisticRegression(solver="liblinear", random_state=42).fit(X_tr3, y_tr)
y_pred_log_tr3 = logreg3.predict(X_tr3)
```

```

y_pred_log_val3 = logreg3.predict(X_val3)

rfc3 = RandomForestClassifier().fit(X_tr3, y_tr)
y_pred_rfc_tr3 = rfc3.predict(X_tr3)
y_pred_rfc_val3 = rfc3.predict(X_val3)

abc3 = AdaBoostClassifier().fit(X_tr3, y_tr)
y_pred_abc_tr3 = abc3.predict(X_tr3)
y_pred_abc_val3 = abc3.predict(X_val3)

gbc3 = GradientBoostingClassifier().fit(X_tr3, y_tr)
y_pred_gbc_tr3 = gbc3.predict(X_tr3)
y_pred_gbc_val3 = gbc3.predict(X_val3)

xgb3 = XGBClassifier().fit(X_tr3, y_tr)
y_pred_xgb_tr3 = xgb3.predict(X_tr3)
y_pred_xgb_val3 = xgb3.predict(X_val3)

```

```

[32]: new_baseline = pd.read_csv("data/charts/new_baseline.csv")
      new_baseline

```

```

[32]:
      Unnamed: 0  Accuracy  F1 Score  Recall  Precision  \
0  Logistic Regression New Baseline  0.807167  0.380952  0.271135  0.640288
1      Random Forest New Baseline  0.811500  0.461685  0.369383  0.615482
2      AdaBoost New Baseline  0.820167  0.448646  0.334349  0.681677
3  Gradient Boosting New Baseline  0.820500  0.466039  0.357959  0.667614
4      XGBoost New Baseline  0.814167  0.455300  0.354912  0.634877

      PR AUC
0  0.498666
1  0.492459
2  0.531694
3  0.542846
4  0.508630

```

**\*\* Observations: \*\***

- Brought it down to top 10 features
- Gradient Boosting has the highest PR AUC Score as well as the highest F1 Score, so we are maximizing Recall and Precision in this model

## 12 Hyperparameter Tuning

```

[ ]: logreg = LogisticRegression()
      params = {'C': [0.001, 0.01, 0.1, 1, 10],
                 'penalty': ['none', 'l1', 'l2', 'elasticnet'],
                 'solver': ['liblinear', 'newton-cg', 'lbfgs', 'sag', 'saga']}

```

```

gslog = GridSearchCV(estimator = logreg,
                     param_grid = params,
                     scoring = 'average_precision',
                     cv = 10,
                     n_jobs = -1).fit(X_tr, y_tr)
y_pred_gslog_tr = gslog.predict(X_tr)
y_pred_gslog_val = gslog.predict(X_val)

rfc = RandomForestClassifier()
params = {'n_estimators': [100, 200, 400, 600, 1000],
          'criterion': ['entropy', 'gini'],
          'max_depth': [5, 8, 15, 25, 30],
          'min_samples_split': [2, 5, 10, 15, 100],
          'min_samples_leaf': [1, 2, 5, 10]}
gsrfc = GridSearchCV(estimator = rfc,
                     param_grid = params,
                     scoring = 'average_precision',
                     cv = 5,
                     n_jobs = -1).fit(X_tr, y_tr)
y_pred_gsrfc_tr = gsrfc.predict(X_tr)
y_pred_gsrfc_val = gsrfc.predict(X_val)

rfc = RandomForestClassifier()
params = {'n_estimators': [100, 200, 400, 600, 1000],
          'criterion': ['entropy', 'gini'],
          'max_depth': [5, 8, 15, 25, 30],
          'min_samples_split': [2, 5, 10, 15, 100],
          'min_samples_leaf': [1, 2, 5, 10]}
gsrfc = GridSearchCV(estimator = rfc,
                     param_grid = params,
                     scoring = 'average_precision',
                     cv = 5,
                     n_jobs = -1).fit(X_tr, y_tr)
y_pred_gsrfc_tr = gsrfc.predict(X_tr)
y_pred_gsrfc_val = gsrfc.predict(X_val)

abc = AdaBoostClassifier()
params = {'n_estimators': [10, 50, 100, 200],
          'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.5]}
gsabc = GridSearchCV(estimator = abc,
                     param_grid = params,
                     n_jobs = -1,
                     cv = 5,
                     scoring = 'average_precision').fit(X_tr, y_tr)
y_pred_gsabc_tr = gsabc.predict(X_tr)

```

```

y_pred_gsabc_val = gsabc.predict(X_val)

gbc = GradientBoostingClassifier()
params = {'n_estimators': [10, 100, 1000],
          'learning_rate': [0.001, 0.01, 0.1],
          'max_depth': [3, 7, 9]}
gsgbc = GridSearchCV(estimator = gbc,
                     param_grid = params,
                     n_jobs = -1,
                     cv = 5,
                     scoring = 'average_precision').fit(X_tr, y_tr)
y_pred_gsgbc_tr = gsgbc.predict(X_tr)
y_pred_gsgbc_val = gsgbc.predict(X_val)

```

**\*\* Best Hyperparameters for each model: \*\***

- Logistic Regression: # Best: 0.522622 using {'C': 1, 'penalty': 'l2', 'solver': 'newton-cg'}
- Random Forest: # Best: 0.565196 using {'criterion': 'gini', 'max\_depth': 8, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 400}
- AdaBoost: # Best: 0.545818 using {'learning\_rate': 0.1, 'n\_estimators': 200}
- Gradient Boosting: # Best: 0.558390 using {'learning\_rate': 0.01, 'max\_depth': 3, 'n\_estimators': 1000}
- XGBoost: # Best: 0.555500 using {'max\_depth': 3, 'min\_child\_weight': 5, 'n\_estimators': 50}

```

[35]: scores3 = pd.read_csv("data/charts/scores3.csv")
scores3

```

```

[35]:
      Unnamed: 0  Accuracy  F1 Score  Recall  \
0  Logistic Regression with GridSearchCV  0.807167  0.380952  0.271135
1      Random Forest with GridSearchCV  0.820667  0.464143  0.354912
2      AdaBoost with GridSearchCV  0.820833  0.443870  0.326733
3  Gradient Boosting with GridSearchCV  0.819667  0.462227  0.354151
4      XGBoost with GridSearchCV  0.818500  0.457939  0.350343

      Precision  PR AUC
0  0.640288  0.498658
1  0.670504  0.545164
2  0.691935  0.524433
3  0.665236  0.543499
4  0.660920  0.539395

```

**\*\* Observations: \*\***

- GridSearchCV was run with the scoring parameter set to find the highest average precision score, which is the PR AUC score.
- Random Forest after hyperparameter tuning has the best accuracy as well as the highest PR AUC score.

- We pickle out Random Forest with GridSearchCV tuned parameters as our best model for now before completing class imbalance methods.

## 13 Class Imbalance Methods

(Will update soon....)

- Initial findings show that there is dramatically improved PR AUC scores without sacrificing any accuracy.
- Accuracy has not improved much, but has not decreased.

## 14 Pickle Out Best Model

```
[ ]: pickle_out = open("../data/best_model.pickle", "wb")
pickle.dump(rfcb, pickle_out)
pickle_out.close()
```

## 15 Holdout Set Prediction

```
[ ]: test = pd.read_csv('../data/testing.csv')
tt = test.drop(["ID"], axis=1)
```

```
[ ]: url = 'https://openexchangerates.org/api/latest.json?
↳app_id=c51b1508fb4145259b1c2fade72a2c04'
response = requests.get(url)
data = response.json()
rate = data['rates']['TWD']
```

```
[ ]: data = [tt]
for d in data:
    d.rename(columns={"PAY_0": "behind1",
                      "PAY_2": "behind2",
                      "PAY_3": "behind3",
                      "PAY_4": "behind4",
                      "PAY_5": "behind5",
                      "PAY_6": "behind6",
                      "BILL_AMT1": "billed1",
                      "BILL_AMT2": "billed2",
                      "BILL_AMT3": "billed3",
                      "BILL_AMT4": "billed4",
                      "BILL_AMT5": "billed5",
                      "BILL_AMT6": "billed6",
                      "PAY_AMT1": "paid1",
                      "PAY_AMT2": "paid2",
                      "PAY_AMT3": "paid3",
                      "PAY_AMT4": "paid4",
```

```

        "PAY_AMT5": "paid5",
        "PAY_AMT6": "paid6",
        "SEX": "gender",
        "EDUCATION": "education",
        "MARRIAGE": "marriage",
        "AGE": "age",
        "LIMIT_BAL": "limit"}, inplace=True)
    d[['limit']] = d[['limit']]/rate
    d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']] =
→d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']].
→divide(rate, axis=1)
    d[['paid1', 'paid2', 'paid3', 'paid4', 'paid5', 'paid6']] = d[['paid1',
→'paid2', 'paid3', 'paid4', 'paid5', 'paid6']].divide(rate, axis=1)
    d['limit'] = d['limit'].apply(lambda x: round(x, 2))
    d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']] =
→d[['billed1', 'billed2', 'billed3', 'billed4', 'billed5', 'billed6']].
→apply(lambda x: round(x, 2))
    d[['paid1', 'paid2', 'paid3', 'paid4', 'paid5', 'paid6']] = d[['paid1',
→'paid2', 'paid3', 'paid4', 'paid5', 'paid6']].apply(lambda x: round(x, 2))
    d.replace({'marriage': {0:3}}, inplace=True)
    d.replace({'education': {5:4, 0:4, 6:4}}, inplace=True)

tt = tt.drop(["Unnamed: 0"], axis=1)

```

```

[ ]: for d in data:
    d['avail6'] = (d.limit - d.billed6) / d.limit
    d['avail5'] = (d.limit - d.billed5) / d.limit
    d['avail4'] = (d.limit - d.billed4) / d.limit
    d['avail3'] = (d.limit - d.billed3) / d.limit
    d['avail2'] = (d.limit - d.billed2) / d.limit
    d['avail1'] = (d.limit - d.billed1) / d.limit
    d['avg_av'] = (d.avail1 + d.avail2 + d.avail3 + d.avail4 + d.avail5 + d.
→avail6) / 6

def delayed_payment(d):
    if (d.behind1 > 0) or (d.behind2 > 0) or (d.behind3 > 0) or (d.behind4 > 0)
→or (d.behind5 > 0) or (d.behind6 > 0):
        return 1
    else:
        return 0
for d in data:
    d['delayed'] = d.apply(delayed_payment, axis=1)

def total_months_with_delayed_payments(d):
    count = 0
    if (d.behind1 > 0):

```

```

        count += 1
    if (d.behind2 > 0):
        count += 1
    if (d.behind3 > 0):
        count += 1
    if (d.behind4 > 0):
        count += 1
    if (d.behind5 > 0):
        count += 1
    if (d.behind6 > 0):
        count += 1
    return count
for d in data:
    d['latemths'] = d.apply(total_months_with_delayed_payments, axis=1)

```

```
[ ]: X_tt = tt[['limit', 'behind1', 'paid2', 'delayed', 'latemths', 'age',
    ↪ 'behind2', 'billed1', 'avg_av', 'avail1']]

```

```
[ ]: pickle_in = open("../data/best_model.pickle", "rb")
    model = pickle.load(pickle_in)

```

```
[ ]: y_pred_tt = model.predict(X_tt)

```

```
[ ]: pickle_out = open("../data/final_prediction.pickle", "wb")
    pickle.dump(y_pred_tt, pickle_out)
    pickle_out.close()

```

### Analysis:

There was not a significant difference in the vanilla model, model with all the engineered features, and model after using feature selection methods. The initial models were selected for the highest accuracy and PR AUC score.

Some of the engineered features created seemed to have a stronger correlation than the original variables. I have to check for collinearity as some of the variables would overlap in context. I am surprised that the demographic features does not have a greater correlation with default. It would seem useful for companies to be able to identify certain demographic groups that are more prone to defaulting.

The metric I used was the PR AUC score, but with an eye to increasing accuracy and PR AUC score, which is the scoring parameters I used in GridSearchCV for hyperparameter tuning. Hyperparameter tuning improved accuracy to 82% from a baseline of 77%, and the highest PR AUC score at around 54%. My initial analysis of implementing class imbalance methods is that it substantially increases the PR AUC score to almost 90%, but accuracy tops out at 82% on the validation set.